

Importing Libraries

```
In [1]: # !pip install xgboost
```

```
In [2]: # !pip install graphviz
```

```
In [3]: # !pip install lightgbm
```

```
In [1]: import math
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.ensemble import GradientBoostingClassifier
import xgboost as xgb
import lightgbm as lgb

import graphviz

from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.model_selection import train_test_split
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import cross_val_predict
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score, roc_curve, roc_auc_score
from sklearn.feature_selection import mutual_info_classif
from sklearn.preprocessing import LabelEncoder
from imblearn.over_sampling import RandomOverSampler
from collections import Counter
```

Importing The Dataset

```
In [2]: lung_cancer = pd.read_csv('lung_cancer.csv')
lung_cancer
```

```
Out[2]:
```

	GENDER	AGE	SMOKING	YELLOW_FINGERS	ANXIETY	PEER_PRESSURE	CHRONIC DISEASE	FATIGUE	ALLERGY	WHEEZING	ALCOHOL CONSUMING
0	M	69	1	2	2	1	1	2	1	2	2
1	M	74	2	1	1	1	2	2	2	1	1
2	F	59	1	1	1	2	1	2	1	2	1
3	M	63	2	2	2	1	1	1	1	1	2
4	F	63	1	2	1	1	1	1	1	2	1
...
304	F	56	1	1	1	2	2	2	1	1	2
305	M	70	2	1	1	1	1	2	2	2	2
306	M	58	2	1	1	1	1	1	2	2	2
307	M	67	2	1	2	1	1	2	2	1	2
308	M	62	1	1	1	2	1	2	2	2	2

309 rows × 16 columns

```
In [3]: lung_cancer.isnull().sum()
```

```
Out[3]: GENDER      0
        AGE        0
        SMOKING    0
        YELLOW_FINGERS 0
        ANXIETY    0
        PEER_PRESSURE 0
        CHRONIC_DISEASE 0
        FATIGUE    0
        ALLERGY    0
        WHEEZING   0
        ALCOHOL_CONSUMING 0
        COUGHING   0
        SHORTNESS_OF_BREATH 0
        SWALLOWING_DIFFICULTY 0
        CHEST_PAIN 0
        LUNG_CANCER 0
dtype: int64
```

```
In [4]: le = LabelEncoder()
```

```
In [5]: lung_cancer.GENDER = le.fit_transform(lung_cancer['GENDER'])
        lung_cancer.LUNG_CANCER = le.fit_transform(lung_cancer['LUNG_CANCER'])
```

```
In [6]: x = lung_cancer.drop(['LUNG_CANCER'], axis = 1)
        y = lung_cancer['LUNG_CANCER']
        mutual_info = mutual_info_classif(x,y)
        mutual_info = pd.Series(mutual_info)
        mutual_info.index = x.columns
        mutual_info.sort_values(ascending=False)
```

```
Out[6]: ALLERGY      0.071304
        WHEEZING     0.051149
        SWALLOWING_DIFFICULTY 0.045092
        AGE          0.034302
        COUGHING     0.033011
        ALCOHOL_CONSUMING 0.027062
        YELLOW_FINGERS 0.016004
        SMOKING      0.012691
        CHEST_PAIN   0.009974
        PEER_PRESSURE 0.006295
        GENDER       0.000000
        ANXIETY      0.000000
        CHRONIC_DISEASE 0.000000
        FATIGUE      0.000000
        SHORTNESS_OF_BREATH 0.000000
dtype: float64
```

```
In [7]: lung_cancer.mean()
```

```
Out[7]: GENDER      0.524272
        AGE        62.673139
        SMOKING     1.563107
        YELLOW_FINGERS 1.569579
        ANXIETY     1.498382
        PEER_PRESSURE 1.501618
        CHRONIC_DISEASE 1.504854
        FATIGUE     1.673139
        ALLERGY     1.556634
        WHEEZING    1.556634
        ALCOHOL_CONSUMING 1.556634
        COUGHING    1.579288
        SHORTNESS_OF_BREATH 1.640777
        SWALLOWING_DIFFICULTY 1.469256
        CHEST_PAIN  1.556634
        LUNG_CANCER 0.873786
dtype: float64
```

```
In [8]: lung_cancer.min()
```

```
Out[8]: GENDER      0
        AGE        21
        SMOKING     1
        YELLOW_FINGERS 1
        ANXIETY     1
        PEER_PRESSURE 1
        CHRONIC_DISEASE 1
        FATIGUE     1
        ALLERGY     1
        WHEEZING    1
        ALCOHOL_CONSUMING 1
        COUGHING    1
        SHORTNESS_OF_BREATH 1
        SWALLOWING_DIFFICULTY 1
        CHEST_PAIN  1
        LUNG_CANCER 0
dtype: int64
```

```
In [9]: lung_cancer.max()
```


```
Out[9]: GENDER      1
        AGE        87
        SMOKING     2
        YELLOW_FINGERS 2
        ANXIETY     2
        PEER_PRESSURE 2
        CHRONIC_DISEASE 2
        FATIGUE     2
        ALLERGY     2
        WHEEZING    2
        ALCOHOL_CONSUMING 2
        COUGHING    2
        SHORTNESS OF BREATH 2
        SWALLOWING DIFFICULTY 2
        CHEST PAIN  2
        LUNG_CANCER 1
        dtype: int64
```

```
In [10]: lung_cancer.std()
```

```
Out[10]: GENDER      0.500221
        AGE        8.210301
        SMOKING     0.496806
        YELLOW_FINGERS 0.495938
        ANXIETY     0.500808
        PEER_PRESSURE 0.500808
        CHRONIC_DISEASE 0.500787
        FATIGUE     0.469827
        ALLERGY     0.497588
        WHEEZING    0.497588
        ALCOHOL_CONSUMING 0.497588
        COUGHING    0.494474
        SHORTNESS OF BREATH 0.480551
        SWALLOWING DIFFICULTY 0.499863
        CHEST PAIN  0.497588
        LUNG_CANCER 0.332629
        dtype: float64
```


Visualizing The Data

```
In [11]: plt.figure(figsize=(6,6))
names = ["Yes", "No"]
count = [(lung_cancer.LUNG_CANCER.values == 1).sum(), (lung_cancer.LUNG_CANCER.values == 0).sum()]
plt.bar(names, count, color = ["Red", "Green"])
plt.title('Lung Cancer', color = 'blue', fontsize= 25)
plt.xlabel('Lung Cancer', fontsize= 18)
plt.ylabel('Count', fontsize= 18)
plt.ylim(0,290)
for i in range(len(names)):
    plt.text(i, count[i], count[i], ha='center', va='bottom', fontsize=16)
plt.show()
```



```
In [12]: plt.figure(figsize = (20, 10))
sns.heatmap(lung_cancer.corr(), annot = True, cmap="plasma")
```

```
Out[12]: <AxesSubplot:>
```




```
In [13]: os = RandomOverSampler(0.7)
x_os, y_os = os.fit_resample(x,y)
print("Before fit {}".format(Counter(y)))
print("After fit {}".format(Counter(y_os)))
```

Before fit Counter({1: 270, 0: 39})

After fit Counter({1: 270, 0: 189})

C:\Users\nitis\anaconda3\lib\site-packages\imblearn\utils_validation.py:586: FutureWarning: Pass sampling_strategy=0.7 as keyword args. From version 0.9 passing these as positional arguments will result in an error warnings.warn()

```
In [14]: plt.figure(figsize=(6,6))
names = ["Yes", "No"]
count = [270, 216]
plt.bar(names, count, color = ["Red", "Green"])
plt.title('Lung Cancer', color = 'blue', fontsize= 25)
plt.xlabel('Lung Cancer', fontsize= 18)
plt.ylabel('Count', fontsize= 18)
plt.ylim(0,290)
for i in range(len(names)):
    plt.text(i, count[i], count[i], ha='center', va='bottom', fontsize=16)
plt.show()
```



```
In [15]: scaler = StandardScaler()
```

```
scaler.fit(x_os)
scaled_x = scaler.transform(x_os)
```

```
In [16]: pca = PCA(n_components=7)
pca.fit(scaled_x)
pca_x = pca.transform(scaled_x)
pca_x.shape
```

```
Out[16]: (459, 7)
```

Models Training

```
In [17]: x_train, x_test, y_train, y_test = train_test_split(pca_x,y_os,test_size = 0.35)

GBM = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0, max_depth=1, random_state=45).fit(x_train,
y_train)

XGB = xgb.XGBClassifier(objective="binary:logistic", random_state=45, eval_metric="auc", n_estimators=50).fit(x_train,
y_train)

LGBM = lgb.LGBMClassifier(num_leaves=31, learning_rate=1, n_estimators=100, random_state=45).fit(x_train, y_train)
```

Evaluating Models


```
In [18]: GBM_pre = GBM.predict(x_test)
GBM_sc = GBM.score(x_test, y_test) * 100
GBM_sc = "{:.2f}".format(GBM_sc)

XGB_pre = XGB.predict(x_test)
XGB_sc = XGB.score(x_test, y_test) * 100
XGB_sc = "{:.2f}".format(XGB_sc)

LGBM_pre = LGBM.predict(x_test)
LGBM_sc = LGBM.score(x_test, y_test) * 100
LGBM_sc = "{:.2f}".format(LGBM_sc)
```

Accuracy


```
In [19]: algo = ['GBM', 'XGB', 'LGBM']
acc = [math.floor(int(float(GBM_sc))),math.floor(int(float(XGB_sc))), math.floor(int(float(LGBM_sc)))]
pert = [GBM_sc, XGB_sc, LGBM_sc]
colors = ['blue','magenta', 'green']
plt.bar(algo, acc,color=colors, edgecolor=colors)
plt.title('Accuracy', color = 'coral', fontsize= 23)
plt.xlabel('Classifier Name', fontsize= 18)
plt.ylabel('Accuracy Percentage', fontsize= 18)
plt.ylim(70,102)
for i in range(len(algo)):
    plt.text(i, acc[i], pert[i], ha='center', va='bottom', fontsize=16)
plt.show()
```




```
In [20]: def print_confusion_matrix(confusion_matrix, class_names, figsize = (7,4), fontsize=14, title = "Confusion Matr
df_cm = pd.DataFrame(
    confusion_matrix, index=class_names, columns=class_names,
)
fig = plt.figure(figsize=figsize)
try:
    heatmap = sns.heatmap(df_cm, annot=True, fmt="d")
except ValueError:
    raise ValueError("Confusion matrix values must be integers.")
heatmap.yaxis.set_ticklabels(heatmap.yaxis.get_ticklabels(), rotation=90, ha='right', fontsize=fontsize)
heatmap.xaxis.set_ticklabels(heatmap.xaxis.get_ticklabels(), rotation=0, ha='center', fontsize=fontsize)
plt.ylabel('Actual', fontsize=18, color='blue')
```

```
plt.xlabel('Prediction', fontsize=18, color='blue')
plt.title(title, fontsize=22, color='blue')
```


```
In [21]: cm = confusion_matrix(y_test, GBM_pre)
print_confusion_matrix(cm,["0", "1"], title="GradientBoostingClassifier (GBM)")
```



```
In [22]: cm = confusion_matrix(y_test, XGB_pre)
print_confusion_matrix(cm,["0", "1"], title="XGBoost (XGB)")
```



```
In [23]: cm = confusion_matrix(y_test, LGBM_pre)
print_confusion_matrix(cm,["0", "1"], title="LightGradientBoostingClassifier (LGBM)")
```



```
In [24]: pd.DataFrame(classification_report(y_test, GBM_pre, output_dict=True))
```

```
Out[24]:
```

	0	1	accuracy	macro avg	weighted avg
precision	0.840000	1.000000	0.925466	0.920000	0.937391
recall	1.000000	0.877551	0.925466	0.938776	0.925466
f1-score	0.913043	0.934783	0.925466	0.923913	0.926276
support	63.000000	98.000000	0.925466	161.000000	161.000000

```
In [25]: pd.DataFrame(classification_report(y_test, XGB_pre, output_dict=True))
```

```
Out[25]:
```

	0	1	accuracy	macro avg	weighted avg
precision	0.863014	1.000000	0.937888	0.931507	0.946397
recall	1.000000	0.897959	0.937888	0.948980	0.937888
f1-score	0.926471	0.946237	0.937888	0.936354	0.938502
support	63.000000	98.000000	0.937888	161.000000	161.000000

```
In [26]: pd.DataFrame(classification_report(y_test, LGBM_pre, output_dict=True))
```

```
Out[26]:
```

	0	1	accuracy	macro avg	weighted avg
precision	0.863014	1.000000	0.937888	0.931507	0.946397
recall	1.000000	0.897959	0.937888	0.948980	0.937888
f1-score	0.926471	0.946237	0.937888	0.936354	0.938502
support	63.000000	98.000000	0.937888	161.000000	161.000000

```
In [46]: x_train, x_test, y_train, y_test = train_test_split(x_os,y_os,test_size = 0.35)

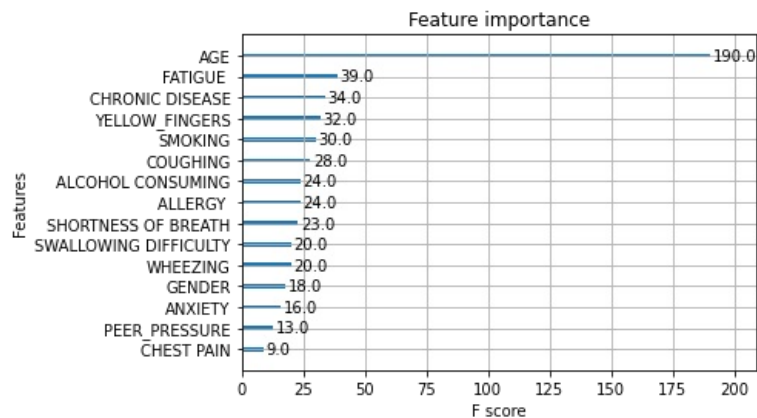
GBM = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0, max_depth=1, random_state=45).fit(x_train, y_train)

XGB = xgb.XGBClassifier(objective="binary:logistic", random_state=45, eval_metric="auc", n_estimators=50).fit(x_train, y_train)

LGBM = lgb.LGBMClassifier(num_leaves=31, learning_rate=1, n_estimators=100, random_state=45).fit(x_train, y_train)
```

```
In [47]: plt.figure(figsize=(10,9))
xgb.plot_importance(XGB)
plt.show()
```

<Figure size 720x648 with 0 Axes>

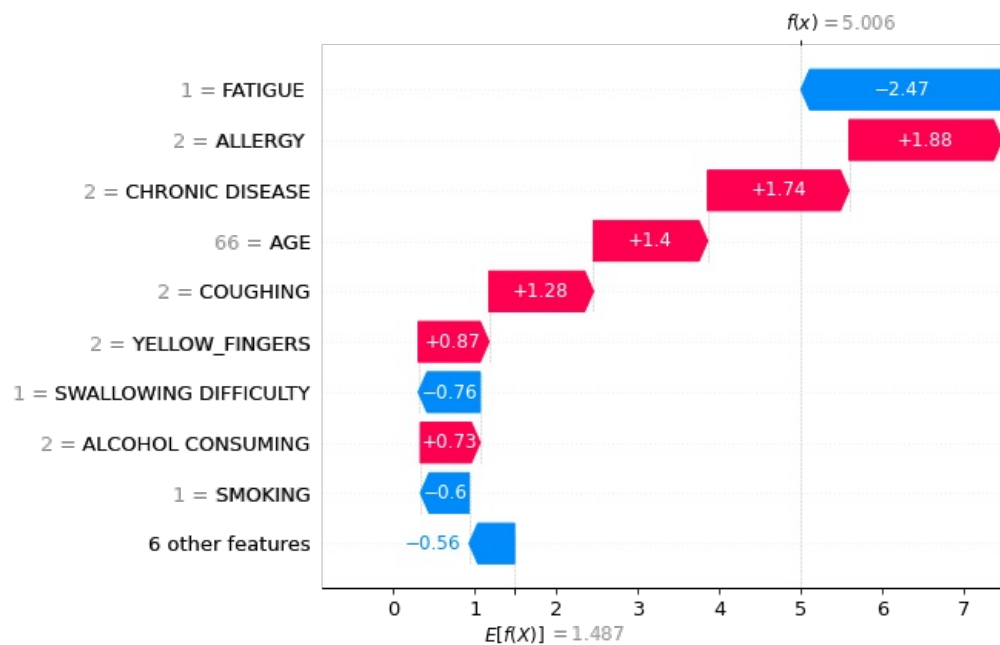


```
In [48]: # !pip install shap
```


```
In [49]: import shap
```

```
In [50]: explainer = shap.Explainer(GBM, x_train)
shap_values = explainer(x_train)
```


```
In [51]: shap.plots.waterfall(shap_values[0])
```



```
In [52]: shap.plots.bar(shap_values)
```

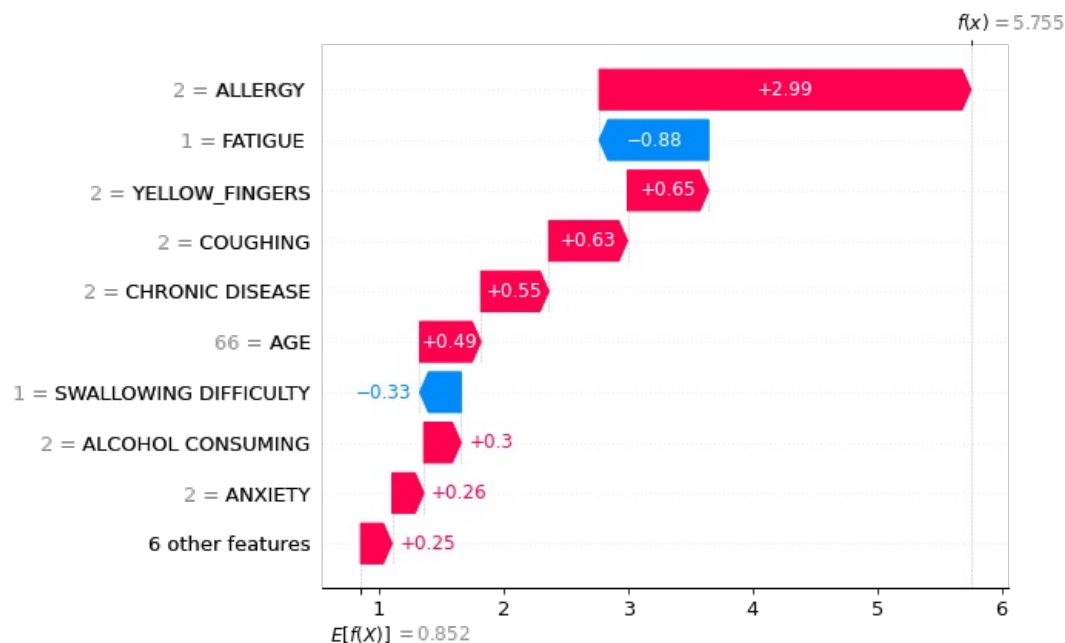


```
In [53]: shap.plots.beeswarm(shap_values)
```

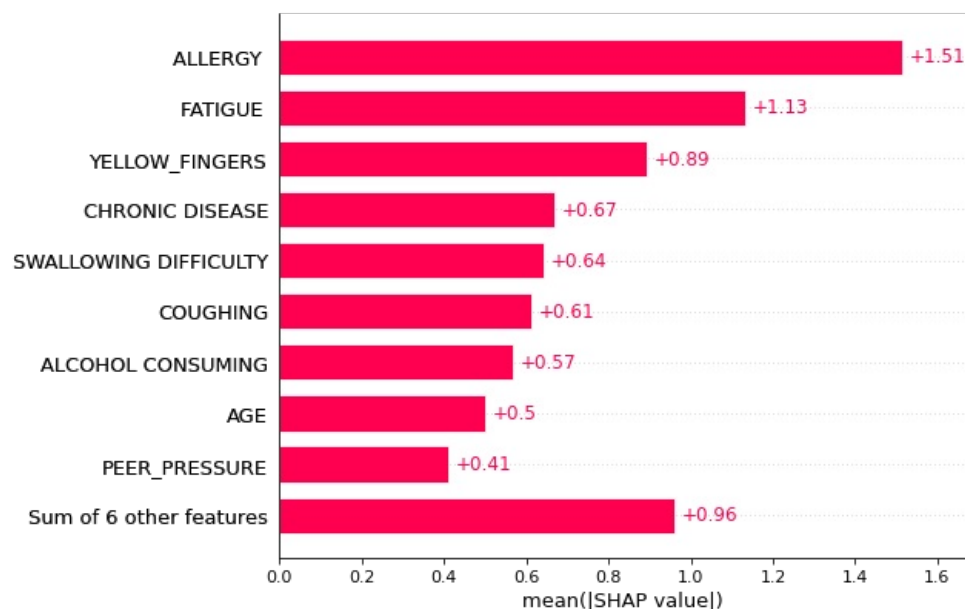


```
In [54]: explainer = shap.Explainer(XGB, x_train)
shap_values = explainer(x_train)
```

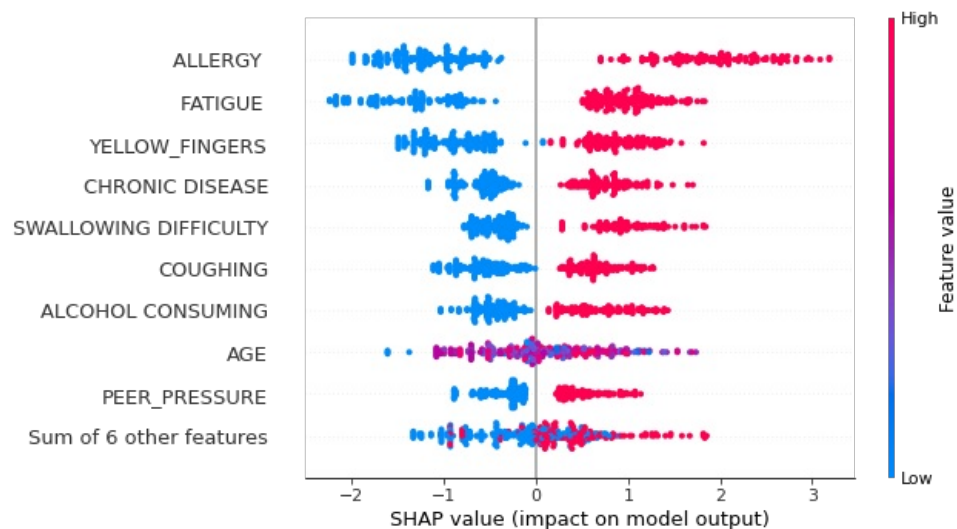
```
In [55]: shap.plots.waterfall(shap_values[0])
```

```
In [56]: shap.plots.bar(shap_values)
```




```
In [57]: shap.plots.beeswarm(shap_values)
```

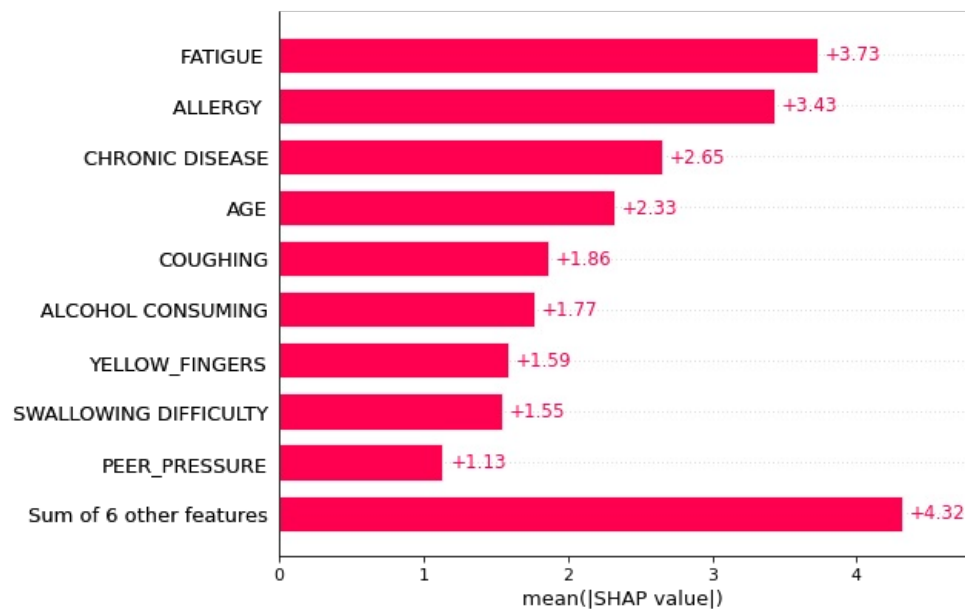


```
In [58]: explainer = shap.Explainer(LGBM, x_train)
shap_values = explainer(x_train)
```

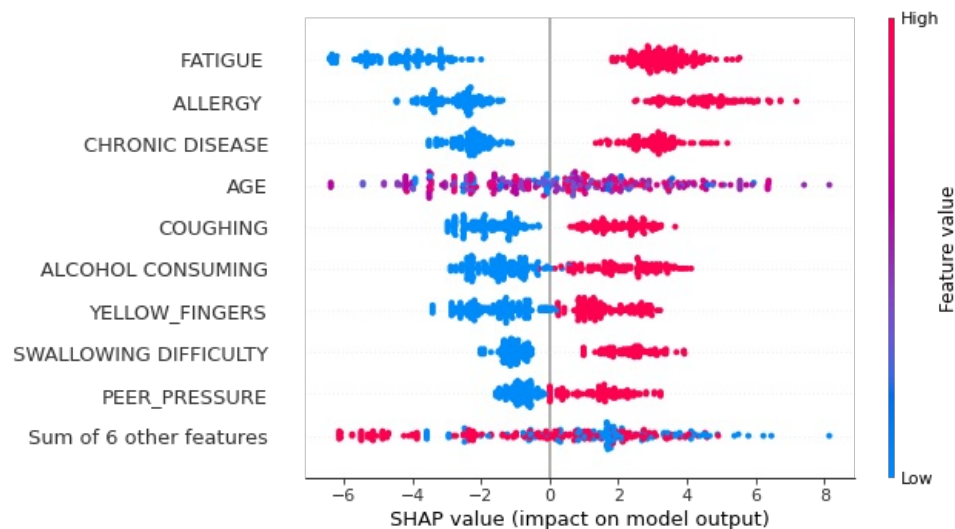
```
In [59]: shap.plots.waterfall(shap_values[0])
```



```
In [60]: shap.plots.bar(shap_values)
```



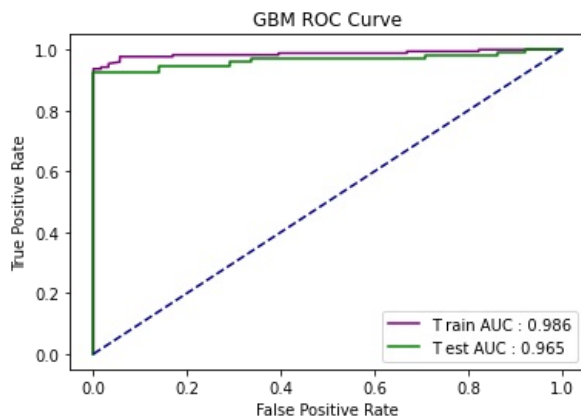
```
In [61]: shap.plots.beeswarm(shap_values)
```



```
In [62]: GBM_probs_tr = GBM.predict_proba(x_train)[: ,1]
auc_tr = roc_auc_score(y_train, GBM_probs_tr)
fpr_tr, tpr_tr, thresholds = roc_curve(y_train, GBM_probs_tr)

GBM_probs_ts = GBM.predict_proba(x_test)[: ,1]
auc_ts = roc_auc_score(y_test, GBM_probs_ts)
fpr_ts, tpr_ts, thresholds = roc_curve(y_test, GBM_probs_ts)
```

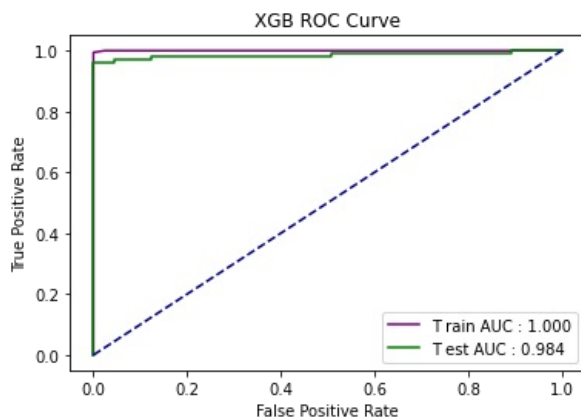
```
plt.plot(fpr_tr, tpr_tr, color='purple', label="T rain AUC : {:.3f}".format(auc_tr))
plt.plot(fpr_ts, tpr_ts, color='green', label="T est AUC : {:.3f}".format(auc_ts))
plt.plot([0,1], [0,1], color='darkblue', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('GBM ROC Curve')
plt.legend()
plt.show()
```



```
In [63]: XGB_probs_tr = XGB.predict_proba(x_train)[: ,1]
auc_tr = roc_auc_score(y_train, XGB_probs_tr)
fpr_tr, tpr_tr, thresholds = roc_curve(y_train, XGB_probs_tr)
```

```
XGB_probs_ts = XGB.predict_proba(x_test)[: ,1]
auc_ts = roc_auc_score(y_test, XGB_probs_ts)
fpr_ts, tpr_ts, thresholds = roc_curve(y_test, XGB_probs_ts)
```


```
plt.plot(fpr_tr, tpr_tr, color='purple', label="T rain AUC : {:.3f}".format(auc_tr))
plt.plot(fpr_ts, tpr_ts, color='green', label="T est AUC : {:.3f}".format(auc_ts))
plt.plot([0,1], [0,1], color='darkblue', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('XGB ROC Curve')
plt.legend()
plt.show()
```



```
In [64]: LGBM_probs_tr = LGBM.predict_proba(x_train)[: ,1]
auc_tr = roc_auc_score(y_train, LGBM_probs_tr)
fpr_tr, tpr_tr, thresholds = roc_curve(y_train, LGBM_probs_tr)
```

```
LGBM_probs_ts = LGBM.predict_proba(x_test)[: ,1]
auc_ts = roc_auc_score(y_test, LGBM_probs_ts)
fpr_ts, tpr_ts, thresholds = roc_curve(y_test, LGBM_probs_ts)
```

```
plt.plot(fpr_tr, tpr_tr, color='purple', label="T rain AUC : {:.3f}".format(auc_tr))
plt.plot(fpr_ts, tpr_ts, color='green', label="T est AUC : {:.3f}".format(auc_ts))
plt.plot([0,1], [0,1], color='darkblue', linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('LGBM ROC Curve')
plt.legend()
plt.show()
```



```
In [65]: # import pickle  
# pickle.dump(GBM, open("GBM_lung_cancer_model.pkl", "wb"))
```

```
In [ ]:
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js